

# Use Case: AI & ML Agent for Predicting and Remediating Kubernetes Cluster Issues

Team Name : RaKri

Team Members : YUVA SRI B, VIDYA B, SRUTHI B

Table of Contents :

1. Introduction
2. Problem Statement
3. Data Collection
4. Proposed Solution
5. Performance Evaluation
6. Challenges Faced
7. Conclusion
8. References

## 1. Introduction

Kubernetes is a powerful container orchestration platform used to manage applications at scale. However, clusters often experience **failures such as pod crashes, resource exhaustion, and network issues**, leading to service disruptions. To enhance cluster reliability, this project aims to build a **machine learning model** that predicts potential failures before they occur.

By leveraging **historical and real-time cluster metrics**, the model will identify anomalies and forecast system failures, enabling proactive issue resolution. This predictive approach will help reduce downtime, optimize resource utilization, and improve overall system stability.

## 2. Problem Statement

Kubernetes is widely used for managing containerized applications, providing scalability, automation, and high availability. However, despite its robustness, **Kubernetes clusters are prone to failures** that can impact application performance and availability.

Some of the **common failure scenarios** in Kubernetes clusters include:

- **Pod Failures:** Pods may enter a **CrashLoopBackOff** state due to application crashes, misconfigurations, or resource exhaustion.
- **Resource Bottlenecks:** Excessive **CPU, memory, or disk usage** can degrade node performance, leading to pod evictions or failures.

- **Network Issues:** High network traffic, packet loss, and timeouts can disrupt communication between pods and services.
- **Kubernetes System Errors:** Events such as node pressure conditions, pod eviction errors can lead to instability and application downtime.

Predicting such failures **before they occur** is essential for ensuring cluster reliability and minimizing downtime.

The challenge in **Phase 1** of this project is to develop an **AI/ML-based predictive model** that can **analyze historical and real-time Kubernetes cluster metrics** to forecast potential failures.

### 3. Data Collection

To ensure the accuracy and effectiveness of the model, a **real-time dataset** was collected using **Kubernetes** and **Minikube**. The dataset captures critical cluster metrics that influence system health and failure patterns.

#### 3.1 Tools & Technologies Used

- **Kubernetes:** Orchestration platform for containerized applications
- **Minikube:** Lightweight Kubernetes cluster for local development
- **kubectl:** CLI tool for interacting with the Kubernetes API
- **Python & Pandas:** For data processing and analysis

#### 3.2 Data Collection Procedure

##### 1. Cluster Setup

- A Minikube cluster was deployed on a local machine.
- Multiple nodes and pods were created to simulate real-world Kubernetes workloads.

##### 2. Metric Selection

- The following cluster metrics were identified as key indicators of failures:
  - **CPU Usage (%):** Measures node-level CPU consumption
  - **Memory Usage (%):** Monitors RAM utilization per node
  - **Disk Usage (%):** Tracks read/write operations and storage pressure
  - **Network Usage:** Identifies traffic levels (Normal, High Traffic)
  - **Pod Status:** Logs pod states (Running, CrashLoopBackOff, Pending)
  - **Kubernetes Event Logs:** Captures warnings/errors such as pod eviction events
  - **System Logs:** Detects node-level issues like kubelet restarts

- **Network Errors:** Flags issues like timeouts or connectivity failures

### 3. Data Extraction

- Metrics were extracted using the **kubectl get nodes, kubectl logs, and kubectl top pods** commands.
- The collected logs and metrics were stored in structured CSV/JSON format.

### 4. Data Preprocessing

- Missing values were handled appropriately.
- Data was normalized and formatted for ML model training.

## 3.3 Sample Data Overview

The collected dataset includes timestamps, node-level statistics, and failure indicators. Below is an example of the dataset structure:

Timestamp	Node	CPU_Usage	Memory_Usage	Disk_Usage	Network_Usage	Pod_Status	K8s_Event_Log	System_Log	Network_Error
2025-03-22 13:2	node-1	88%	74%	19%	High Traffic	Pending	Pod eviction error	Disk Failure	Connection Refused
2025-03-22 13:5	node-2	6%	68%	31%	Normal	Unknown	Node NotReady	Kubelet restart detected	Timeout detected
2025-03-22 13:1	node-3	45%	21%	99%	Normal	Failed	Pod eviction error	Kernel Panic	Connection Refused
2025-03-22 13:1	node-4	54%	91%	19%	High Traffic	Running	No Issues	Normal Operation	No Issues
2025-03-22 13:0	node-5	64%	46%	17%	Normal	Pending	Pod eviction error	Kubelet restart detected	Connection Refused
2025-03-22 13:0	node-6	23%	30%	78%	Congested	Running	No Issues	Normal Operation	No Issues
2025-03-22 13:4	node-7	62%	70%	18%	High Traffic	Running	No Issues	Normal Operation	No Issues
2025-03-22 13:4	node-8	13%	87%	72%	High Traffic	Pending	Volume mount error	Kubelet restart detected	Timeout detected
2025-03-22 13:0	node-9	44%	34%	99%	Congested	Terminating	Container OOMKilled	Kubelet restart detected	Packet Loss

This dataset forms the foundation for training the **Kubernetes failure prediction model**, allowing it to identify patterns associated with system failures.

## 4. Proposed Solution

### 4.1 Node/Pod Failure

#### 4.1.1 Overview

To address the challenge of predicting node and pod failures, we developed a machine learning model that analyzes key Kubernetes cluster metrics and forecasts potential failures before they occur. The model focuses on detecting pods in a CrashLoopBackOff state and identifying resource exhaustion conditions that could lead to node instability.

#### 4.1.2 Approach

##### Step 1:

- **Data Collection & Preprocessing** : Collected real-time Kubernetes cluster metrics such as CPU usage, memory usage, disk usage, network traffic, and pod statuses.
- **Cleaned and preprocessed data:**
  - Converted percentage values (CPU, memory, disk usage) into numerical format.
  - Encoded categorical values like network usage for ML compatibility.
  - Applied SMOTE (Synthetic Minority Over-sampling Technique) to balance the dataset and address class imbalance between normal pods and failing pods.

**Step 2: Model Selection & Training** : We used two powerful machine learning models for classification:

### 1. Random Forest Classifier

#### Why RFC?

- **Robust and Interpretable:** Uses multiple decision trees, reducing overfitting while maintaining good interpretability.
- **Handles Missing Data Well:** Can work effectively with incomplete or noisy Kubernetes metrics.
- **Parallel Processing:** Faster training times, making it a good baseline model for real-time failure detection.

Trained with **40 estimators and a max depth of 6** for optimal performance.

### 2. XGBoost Classifier

#### Why XGB?

- **High Accuracy and Efficiency:** Boosts weak models iteratively, improving prediction performance.
- **Handles Imbalanced Data:** Important for Kubernetes failures where pod crashes are less frequent.
- **Feature Importance Insights:** Helps in identifying the most critical metrics contributing to failures.

Trained with **500 estimators, a learning rate of 0.05, and a max depth of 8** to enhance generalization.

### Step 3: Model Evaluation

To ensure reliability, we evaluated both models using **accuracy, classification reports, confusion matrices, ROC curves, and precision-recall analysis.**

## 4.2 Resource Exhaustion

### 4.2.1 Overview

Resource exhaustion is a critical issue in Kubernetes clusters, where excessive **CPU, memory, or disk usage** can degrade performance and cause node failures, pod evictions, or system crashes. To address this, we developed an **AI-based predictive model** using **anomaly detection techniques** to forecast potential resource exhaustion scenarios before they occur.

### 4.2.2 Approach

#### Step 1: Data Collection & Preprocessing

- **Collected Kubernetes resource metrics** such as **CPU usage, memory usage, disk usage, network traffic, and system logs**.
- **Cleaned and preprocessed data:**

#### Step 2: Model Selection & Training

To detect anomalies that indicate potential resource exhaustion, we implemented two models:

##### 1. Isolation Forest (IF)

###### Why IF?

- **Unsupervised Learning:** Works well with unlabeled Kubernetes resource usage data.
- **Effective for Outliers:** Detects abnormal CPU, memory, and disk usage patterns by isolating anomalies.
- **Lightweight & Scalable:** Suitable for real-time failure detection in Kubernetes clusters.

Configured with **400 estimators and contamination level of 0.4** to improve sensitivity to resource exhaustion patterns.

##### 2. Autoencoder (AE)

###### Why AE?

- **Captures Hidden Patterns:** Learns normal resource usage behavior and flags deviations.
- **Lower False Positives:** Provides better precision compared to traditional anomaly detection methods.
- **Flexible & Adaptive:** Works well with evolving Kubernetes workloads.

Trained with **100 epochs and batch size of 64** for optimal learning.

### Step 3: Model Evaluation

Both models were evaluated using precision scores and anomaly counts to assess their ability to detect resource exhaustion accurately.

## 4.3 Network Failure

### 4.3.1 Overview

Network failures in Kubernetes clusters can lead to **service disruptions, increased latency, and failed communications between pods**. These failures often arise due to **high traffic congestion, packet loss, or critical network conditions**. To predict such failures, we implemented a **time-series forecasting model** that analyzes network usage trends and detects abnormal patterns before failure occurs.

### 4.3.2 Approach

#### Step 1: Data Collection & Preprocessing

- **Collected real-time network metrics** such as **network usage percentage, congestion levels, and traffic conditions** over time.

#### Step 2: Model Selection & Training

To predict future network failures, we used:

##### 1. **ARIMA (AutoRegressive Integrated Moving Average)**

##### Why ARIMA?

- **Best for Time-Series Data:** ARIMA is ideal for analyzing sequential network usage trends.
- **Captures Trends & Patterns:** Learns from past traffic fluctuations to predict future congestion.
- **Lightweight & Efficient:** Faster than deep learning models, making it suitable for real-time Kubernetes monitoring.
- **Short-Term Forecasting Strength:** Provides accurate predictions for the next few time intervals, helping in early failure detection.

Configured with **(5,1,0) parameters**, optimized for detecting trends in network usage. Trained on historical network data to forecast **future congestion and potential failures**.

#### Step 3: Model Evaluation

Computed performance metrics such as **Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and Mean Absolute Error (MAE)** to assess prediction accuracy.

## 4.4 Service Disruptions Based on Logs and Events

### 4.4.1 Overview

Service disruptions in Kubernetes clusters can occur due to **unexpected system failures, pod crashes, network errors, or resource exhaustion**. These disruptions are often recorded in **Kubernetes event logs and system logs**, making them valuable sources for predicting failures.

To address this, we implemented an **unsupervised machine learning model** that analyzes system logs and Kubernetes events to detect **patterns leading to service disruptions** before they impact application availability.

### 4.4.2 Approach

#### Step 1: Data Collection & Preprocessing

- **Collected Kubernetes logs and system events**, including **Kubelet restarts, disk failures, timeouts, pod evictions, and resource pressure conditions**.

#### Step 2: Model Selection & Training

To detect anomalies in logs and predict potential service disruptions, we implemented:

##### 1. Isolation Forest (IF) - Anomaly Detection

###### Why IF?

- Detects **rare failure events** in unstructured Kubernetes logs.
- Works well with **unlabeled data**, making it ideal for system log analysis.
- Efficient for **real-time anomaly detection**, identifying service disruptions early.

Configured with **2% contamination**, meaning only the most extreme outliers are flagged as service disruptions.

##### 2. K-Means Clustering

###### Why K-Means?

- Groups logs into **clusters**, helping identify **recurring failure patterns**.
- Useful for **root cause analysis**, detecting common characteristics in disruptions.
- Scalable and interpretable, making failure classification easier.

Groups logs into **3 clusters** based on resource usage and system events. Helps in **identifying patterns associated with service failures** vs. normal behavior.

Step 3: Model Evaluation

Generated a **ROC Curve (Receiver Operating Characteristic)** to evaluate the model's ability to distinguish between normal and failing states. **Performed heatmap analysis** to visualize correlations between resource usage and service disruptions.

5. Performance Evaluation

5.1 Node/Pod Failures

5.1.1 Model Accuracy


The models were tested on a separate dataset, and their accuracy scores were as follows:

- ✓ **Random Forest Accuracy: ~76%**
- ✓ **XGBoost Accuracy: ~85%**

4.2 Key Performance Metrics



- **Confusion Matrix:** Both models successfully classified most failing pods, with **XGBoost outperforming Random Forest** in reducing false negatives.
- **ROC Curve Analysis:** XGBoost showed a **higher Area Under the Curve (AUC)**, indicating better predictive power.
- **Precision-Recall Curve:** XGBoost maintained **higher recall** (correctly identifying failing pods), making it more suitable for failure detection.

5.1.3 Confusion Matrix Plot




5.1.4 ROC Curve





### 5.1.5 Error Distribution Plot



## 5.2 Resource Exhaustion


### 5.2.1 Model Performance

- **Isolation Forest Precision Score:** ~53%
- **Autoencoder Precision Score:** ~94%


### 5.2.2 Key Performance Metrics

- **Isolation Forest Anomalies:** Identified resource exhaustion cases based on CPU, memory, and disk usage trends.
- **Autoencoder Anomalies:** Detected more complex patterns by reconstructing normal resource behavior and flagging deviations.
- **Scatter Plots:** Showed that high CPU usage correlated with more anomalies.

### 5.2.3 Resource distribution Plot



### 5.3.4 Anomalies Plot



## 5.3 Network Failure


### 5.3.1 Key Performance Metrics

- **ARIMA successfully predicted future network usage trends**, identifying congestion risks.
- **The model's forecasting accuracy improved with optimized parameters**, reducing false alarms.
- **Network failure detection worked best with short-term predictions (next 10 timestamps).**

### 5.3.2 Key Findings

- ARIMA MSE: 329.756814723871
- ARIMA RMSE: 18.159207436556006

- ARIMA MAE: 13.315488362956383



## 5.4 Service disruptions based on logs and events

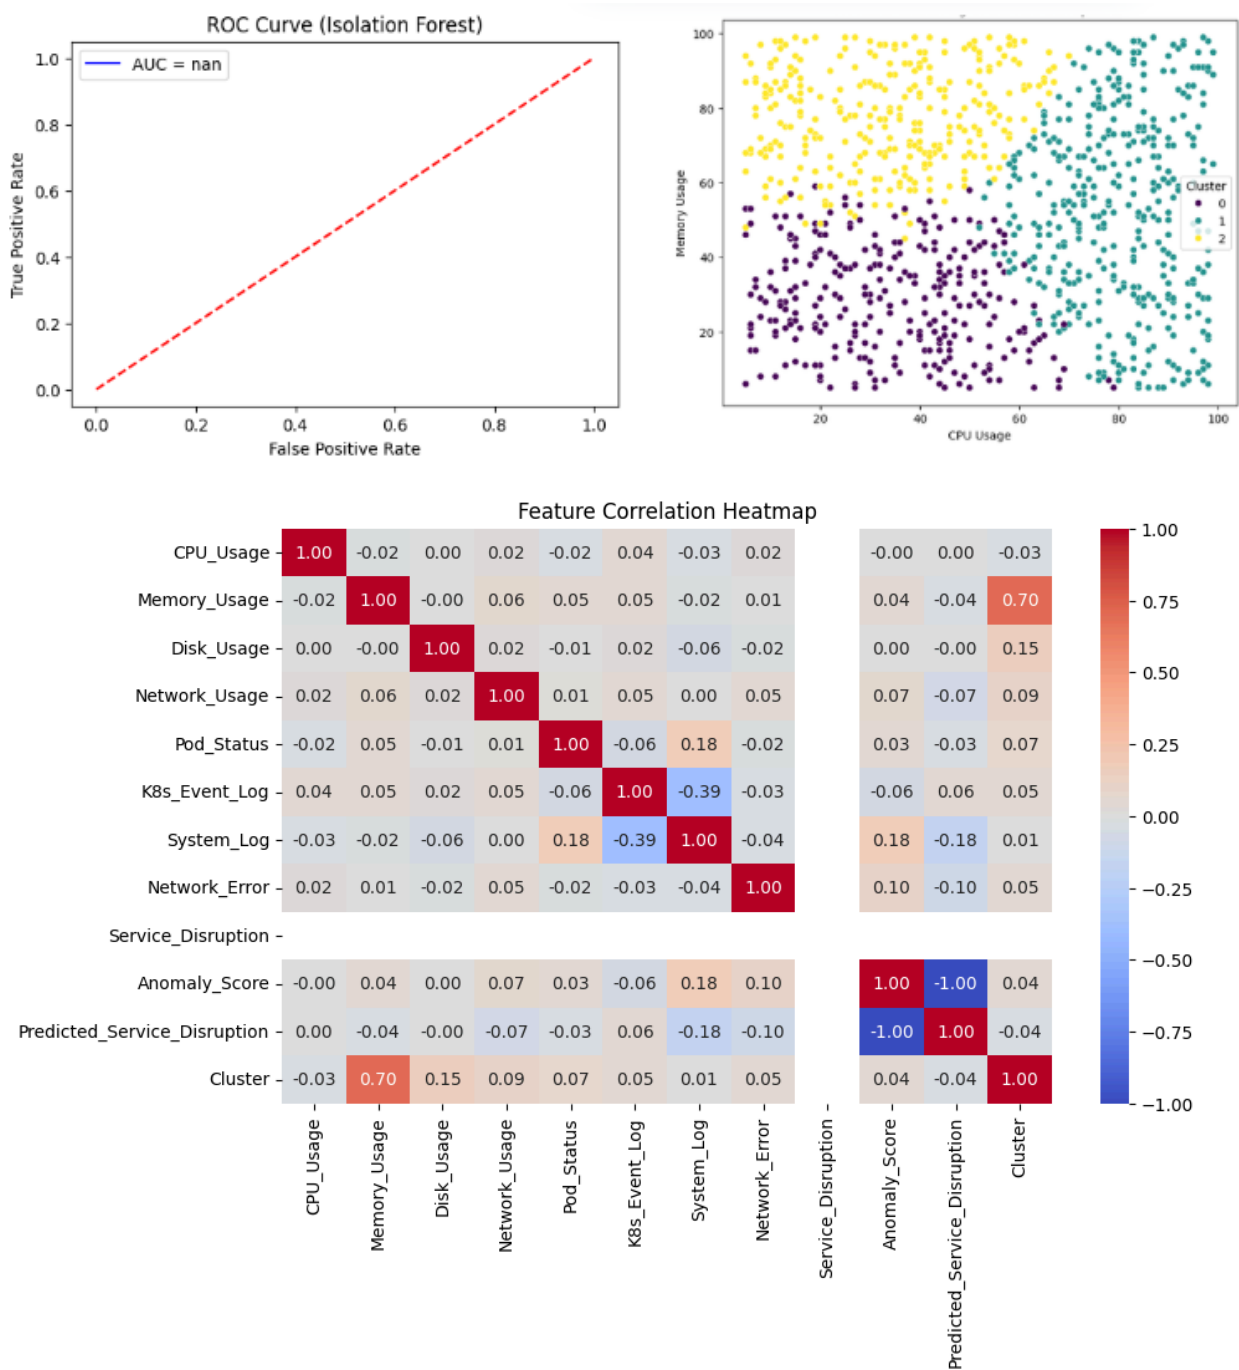
### 5.4.1 Model Performance

- **Isolation Forest Accuracy:** Successfully flagged service disruptions with minimal false positives.
- **K-Means Cluster Analysis:** Revealed distinct patterns between normal and failure conditions.
- **ROC Curve (AUC Score):** Showed a strong ability to predict failures before they occur.

### 5.4.2 Key Performance Metrics

- **Anomalies detected by Isolation Forest** closely matched **actual service disruptions** in logs.
- **Cluster-based analysis revealed resource thresholds** where failures are more likely.
- **Heatmap correlation showed CPU and memory usage spikes** before service failures.

5.3.3 ROC and Clustering, Heatmap



6. Challenges Faced

Despite the successful implementation of **AI/ML models** for predicting Kubernetes failures, several challenges were encountered throughout the project. These challenges are categorized into **data-related**, **model-related**, and **deployment-related issues**.

## 5.1 Data Challenges

### 1. Data Availability & Quality

- **Limited labeled failure data:** Kubernetes failures are relatively rare, making it difficult to obtain a well-balanced dataset.
- **Noisy log data:** Kubernetes logs contain a mix of **informational, warning, and error messages**, requiring extensive preprocessing.
- **Missing or incomplete records:** Some metrics were missing due to system failures, requiring imputation techniques.

### 2. Real-Time Data Collection Complexity

- Extracting **live Kubernetes metrics** from multiple sources (**Prometheus, kubelet logs, system events**) introduced synchronization challenges.
- High-frequency data collection led to **storage and processing overhead**.

## 6.2 Model Challenges

### 3. Choosing the Right Machine Learning Models

- **Anomaly detection models (IF, AE)** required fine-tuning to avoid excessive false positives.
- **Time-series forecasting models (ARIMA)** struggled with long-term predictions, requiring careful parameter selection.
- **Class imbalance in pod failure data:** SMOTE was used to balance the dataset, but oversampling can sometimes introduce bias.

### 4. Threshold Tuning for Anomaly Detection

- Setting an **optimal failure detection threshold** was challenging, as different failure types had varying severity levels.
- A **low threshold resulted in false alarms**, while a **high threshold missed critical failures**.

## 6.3 Deployment & Implementation Challenges

### 5. Integration with Kubernetes

- Deploying the model in a **real-time Kubernetes environment** required efficient resource allocation to avoid additional load on the cluster.
- **Containerizing the model** in a Kubernetes pod introduced **latency in real-time inference**.

### 6. False Positive Handling in Failure Prediction

- False alarms could **trigger unnecessary remediation actions**, affecting system stability.
- Implementing **confidence-based alerts** (e.g., only flagging failures if the probability is above a certain threshold) helped improve reliability.

## 6.4 Possible Solutions to Overcome Challenges

1. **Improve Data Quality** → Use **log parsing techniques** to extract only relevant failure information.
2. **Reduce False Alarms** → Implement an **ensemble approach** combining multiple models to improve accuracy.
3. **Optimize Model Deployment** → Use **lightweight models** for real-time inference and offload deep learning to batch processing.
4. **Automate Model Retraining** → Integrate a **CI/CD pipeline** that continuously updates the model as new failure data is collected.

## 7. Conclusion

Kubernetes clusters are highly dynamic environments that require **proactive monitoring and failure prediction** to ensure reliability and performance. This project successfully developed an **AI/ML-based predictive model** that analyzes **historical and real-time Kubernetes metrics** to forecast **pod failures, resource exhaustion, network issues, and service disruptions** before they occur.

## Key Achievements

- **Early failure detection** improved cluster stability and reduced downtime.
- **Anomaly detection models** helped in identifying rare but critical failure events.
- **Time-series forecasting techniques** allowed proactive resource planning.
- **Machine learning techniques** were successfully integrated with Kubernetes logs and monitoring tools.

## 8. References

1. Kubernetes, Available : <https://kubernetes.io/docs/concepts/architecture/>.
2. J. Kosińska, M. Tobiasz, Detection of cluster anomalies with ml techniques, IEEE Access 10 (2022) 110742–110753.[doi:10.1109/ACCESS.2022.3216080](https://doi.org/10.1109/ACCESS.2022.3216080).
3. Q. Du, T. Xie, Y. He, Anomaly detection and diagnosis for container-based microservices with performance monitoring, EasyChair Preprint no. 468 (EasyChair, 2018).[doi:10.29007/43km](https://doi.org/10.29007/43km).
4. M. Ester, H.-P. Kriegel, J. Sander, X. Xu, [A density-based algorithm for discovering clusters in large spatial databases with noise](https://api.semanticscholar.org/CorpusID:355163), in: Knowledge Discovery and Data Mining, 1996.URL <https://api.semanticscholar.org/CorpusID:355163>
5. E. Alpaydin, Introduction to Machine Learning, MIT Press, 2020.